

Unit 3:

Data Mining

Classification: Basic Concepts, Decision Trees, and Model Evaluation

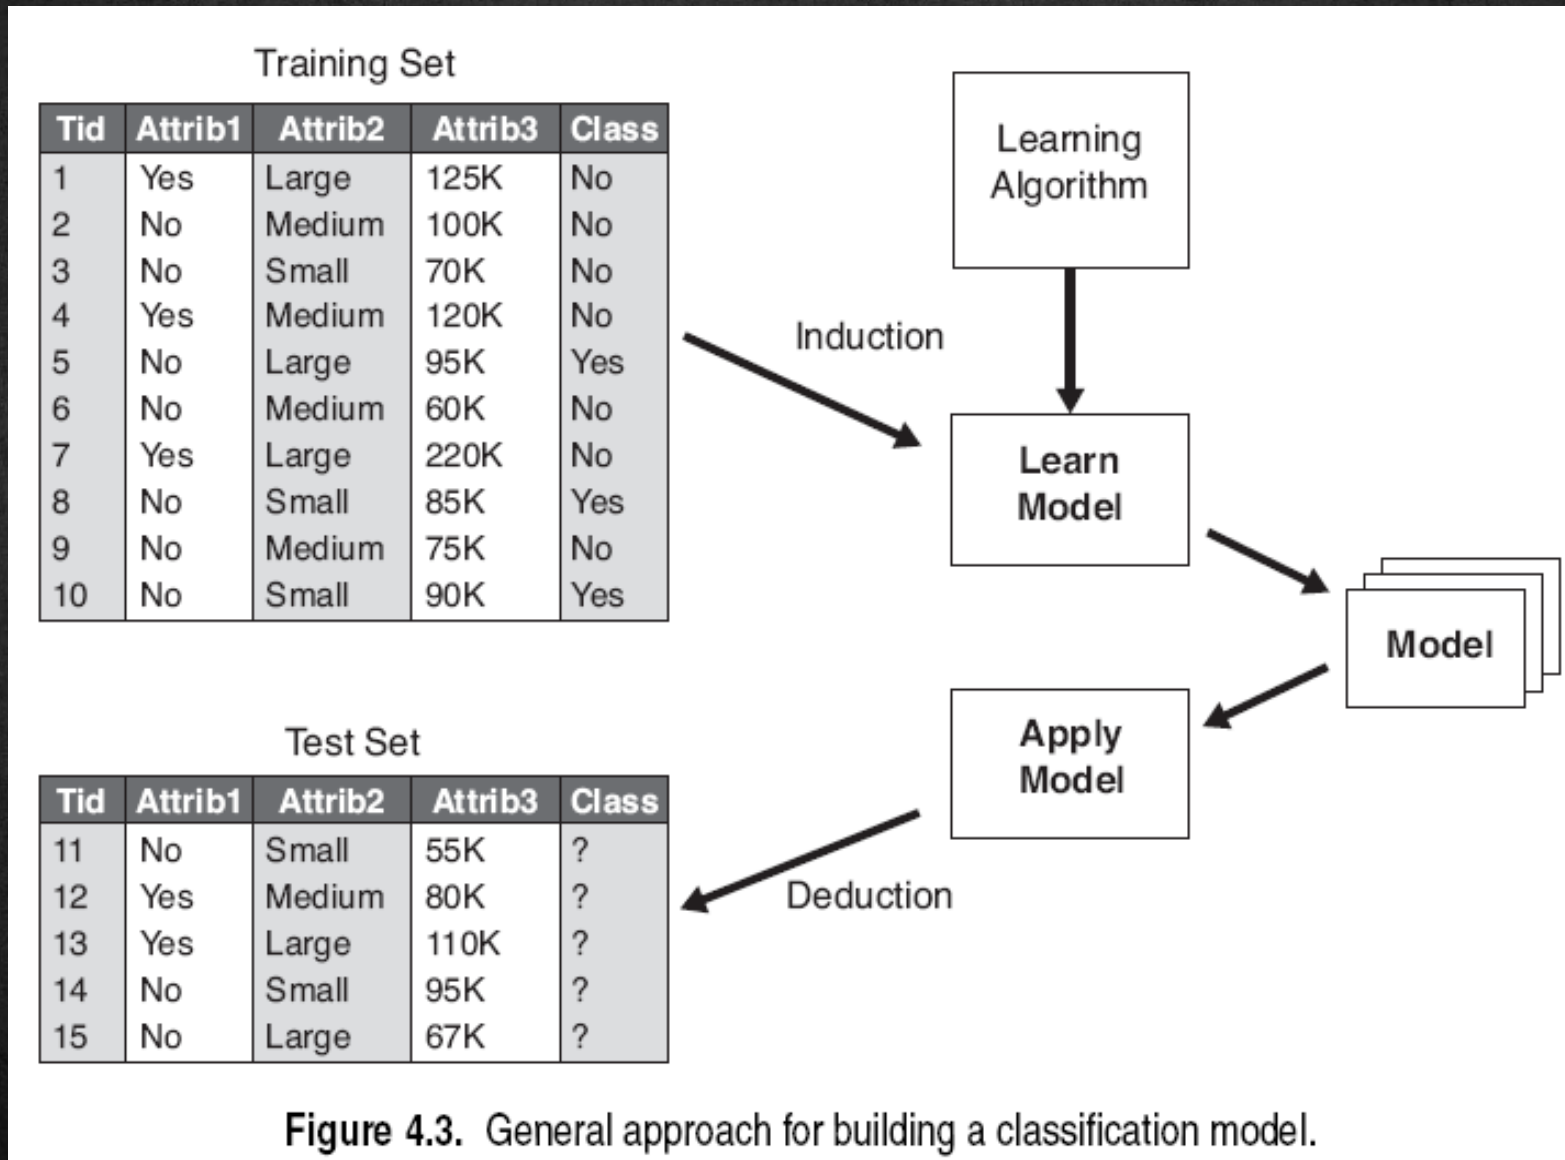
Unit 3

Section 1: Basic Classification

Classification: Definition

- Given a collection of records (training set)
 - Each record contains a set of attributes, one of the attributes is the class.
- Find a model for class attribute as a function of the values of other attributes.
- Goal: previously unseen records should be assigned a class as accurately as possible.
 - A test set is used to determine the accuracy of the model.
 - Usually, the given data set is divided into training and test sets, with training set used to build the model and test set used to validate it.
- Probabilistic vs. non-probabilistic models

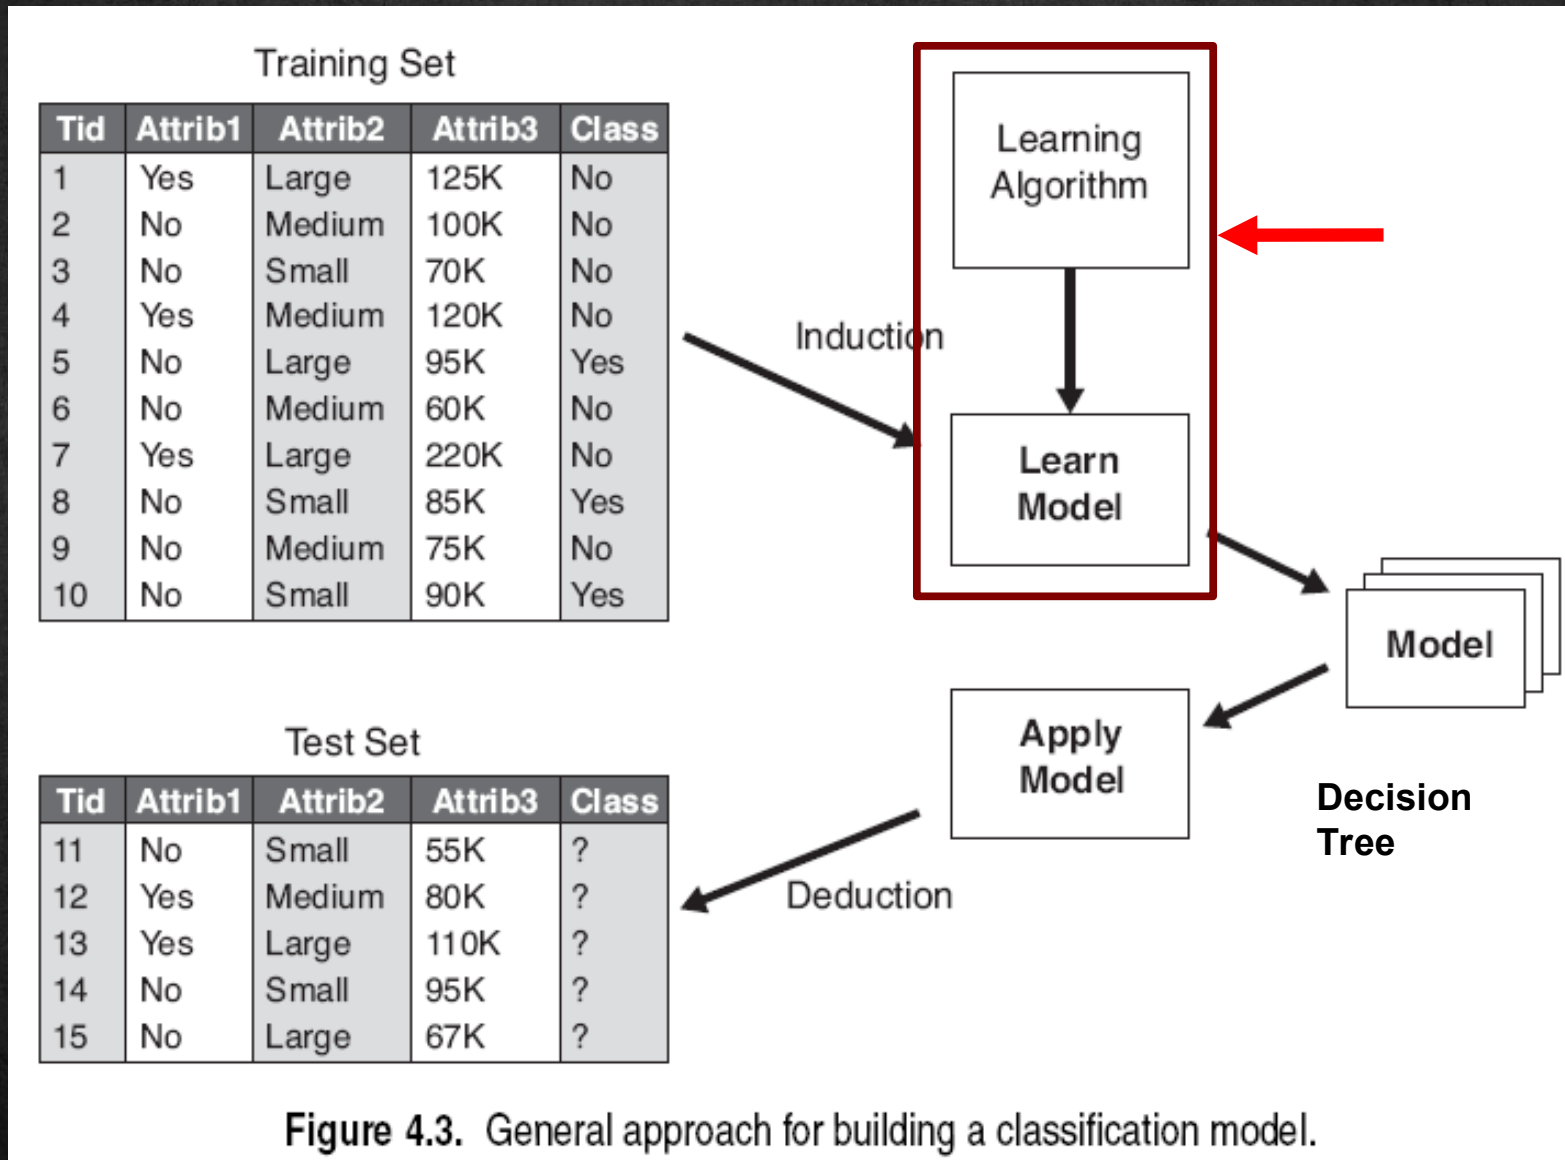
Illustrating Classification Task



Classification Techniques

- **Decision Trees**
- **Rule-based Methods**
- Memory based reasoning
- Neural Networks
- Naïve Bayes and Bayesian Belief Networks
- **Support Vector Machines**
- **Instance-based methods**
- Other methods

Decision Tree Classification Task



Decision Tree Induction


➤ Several Algorithms:

- Hunt's Algorithm: Base of all methods
- CART
- ID3, C4.5
- SLIQ, SPRINT

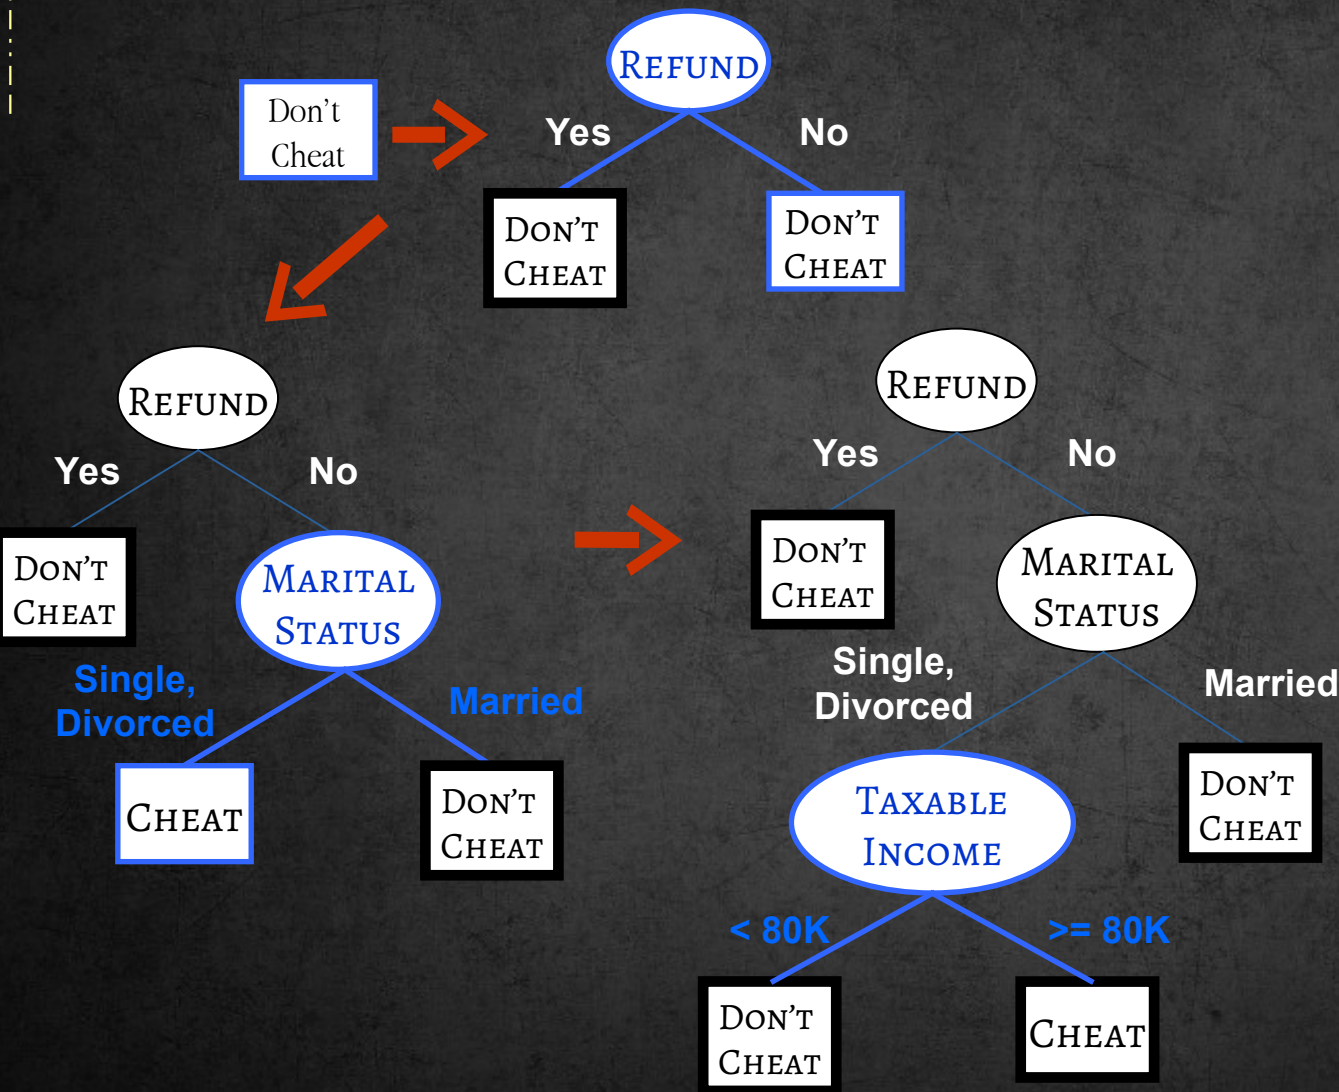
General Structure of Hunt's Algorithm

- Let D_t be the set of training records that reach a node t
- General Procedure:
 - If D_t contains records that belong the same class y_t , then t is a leaf node labeled as y_t
 - If D_t is an empty set, then t is a leaf node labeled by the default class, y_d
 - If D_t contains records that belong to more than one class, use an attribute test to split the data into smaller subsets. Recursively apply the procedure to each subset.

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



Hunt's Algorithm



Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Tree Induction

- Greedy strategy.
 - Split the records based on an attribute test that optimizes certain criterion.
- Issues
 - Determine how to split the records
 - How to specify the attribute test condition?
 - How to determine the best split?
 - Determine when to stop splitting


How to Specify Test Condition?

- Depends on attribute types
 - Nominal
 - Ordinal
 - Continuous


- Depends on number of ways to split
 - 2-way split
 - Multi-way split

Splitting Based on Nominal Attributes


- Multi-way split: Use as many partitions as distinct values.



- Binary split: Divides values into two subsets. Need to find optimal partitioning.




OR




Splitting Based on Ordinal Attributes


- **Multi-way split:** Use as many partitions as distinct values.




- **Binary split:** Divides values into two subsets. Need to find optimal partitioning.



OR



- What about this split?




Splitting Based on Continuous Attributes

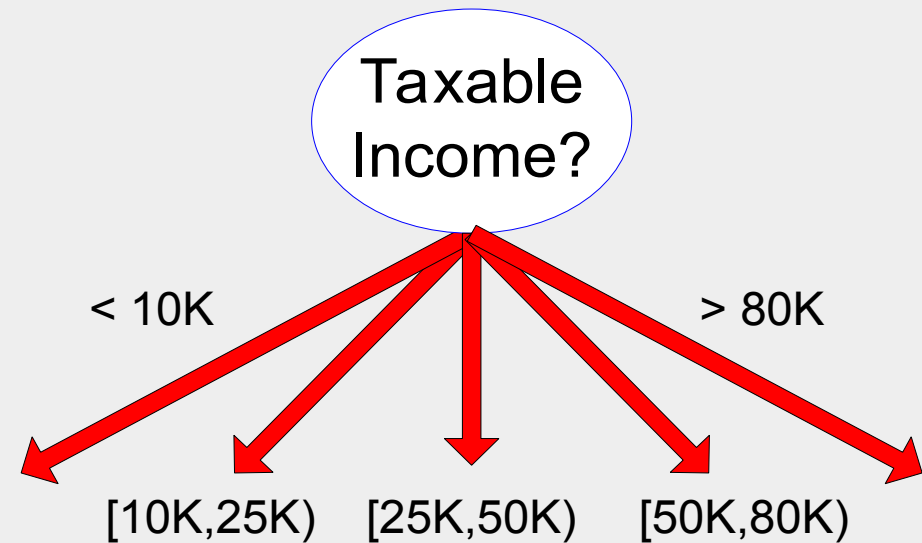
➤ Different ways of handling

- Discretization to form an ordinal categorical attribute
 - Static – discretize once at the beginning
 - Dynamic – ranges can be found by equal interval bucketing, equal frequency bucketing (percentiles), or clustering.
- Binary decision (consider all values): $(A < v)$ or $(A \geq v)$
 - consider all possible splits and finds the best cut
 - can be more compute intensive
 - In some case too many splits

Splitting Based on continuous Attributes



(i) Binary split



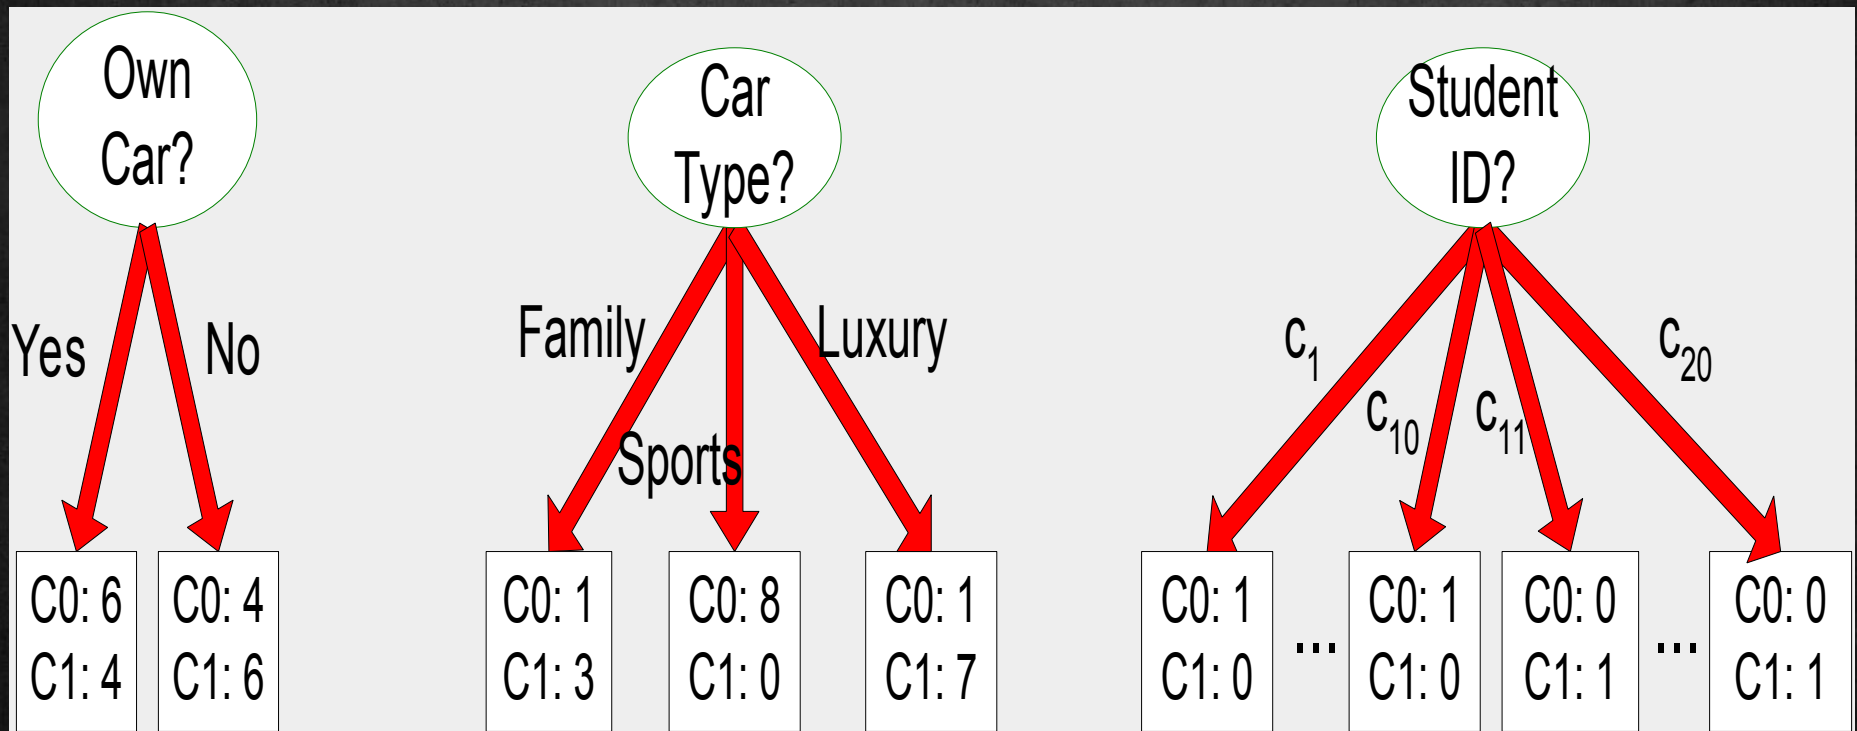
(ii) Multi-way split

Tree Induction

- Greedy strategy.
 - ⦿ Split the records based on an attribute test that optimizes certain criterion.
- Issues
 - ⦿ Determine how to split the records
 - ⦿ How to specify the attribute test condition?
 - ⦿ How to determine the best split?
 - ⦿ Determine when to stop splitting

How to determine the Best Split

Before Splitting: 10 records of class 0, 10 records of class 1



Which test condition is the best?

How to determine the Best Split

- Greedy approach:
 - ◉ Nodes with homogeneous class distribution are preferred
- Need a measure of node impurity:

C0: 5
C1: 5

Non-homogeneous,
High degree of impurity

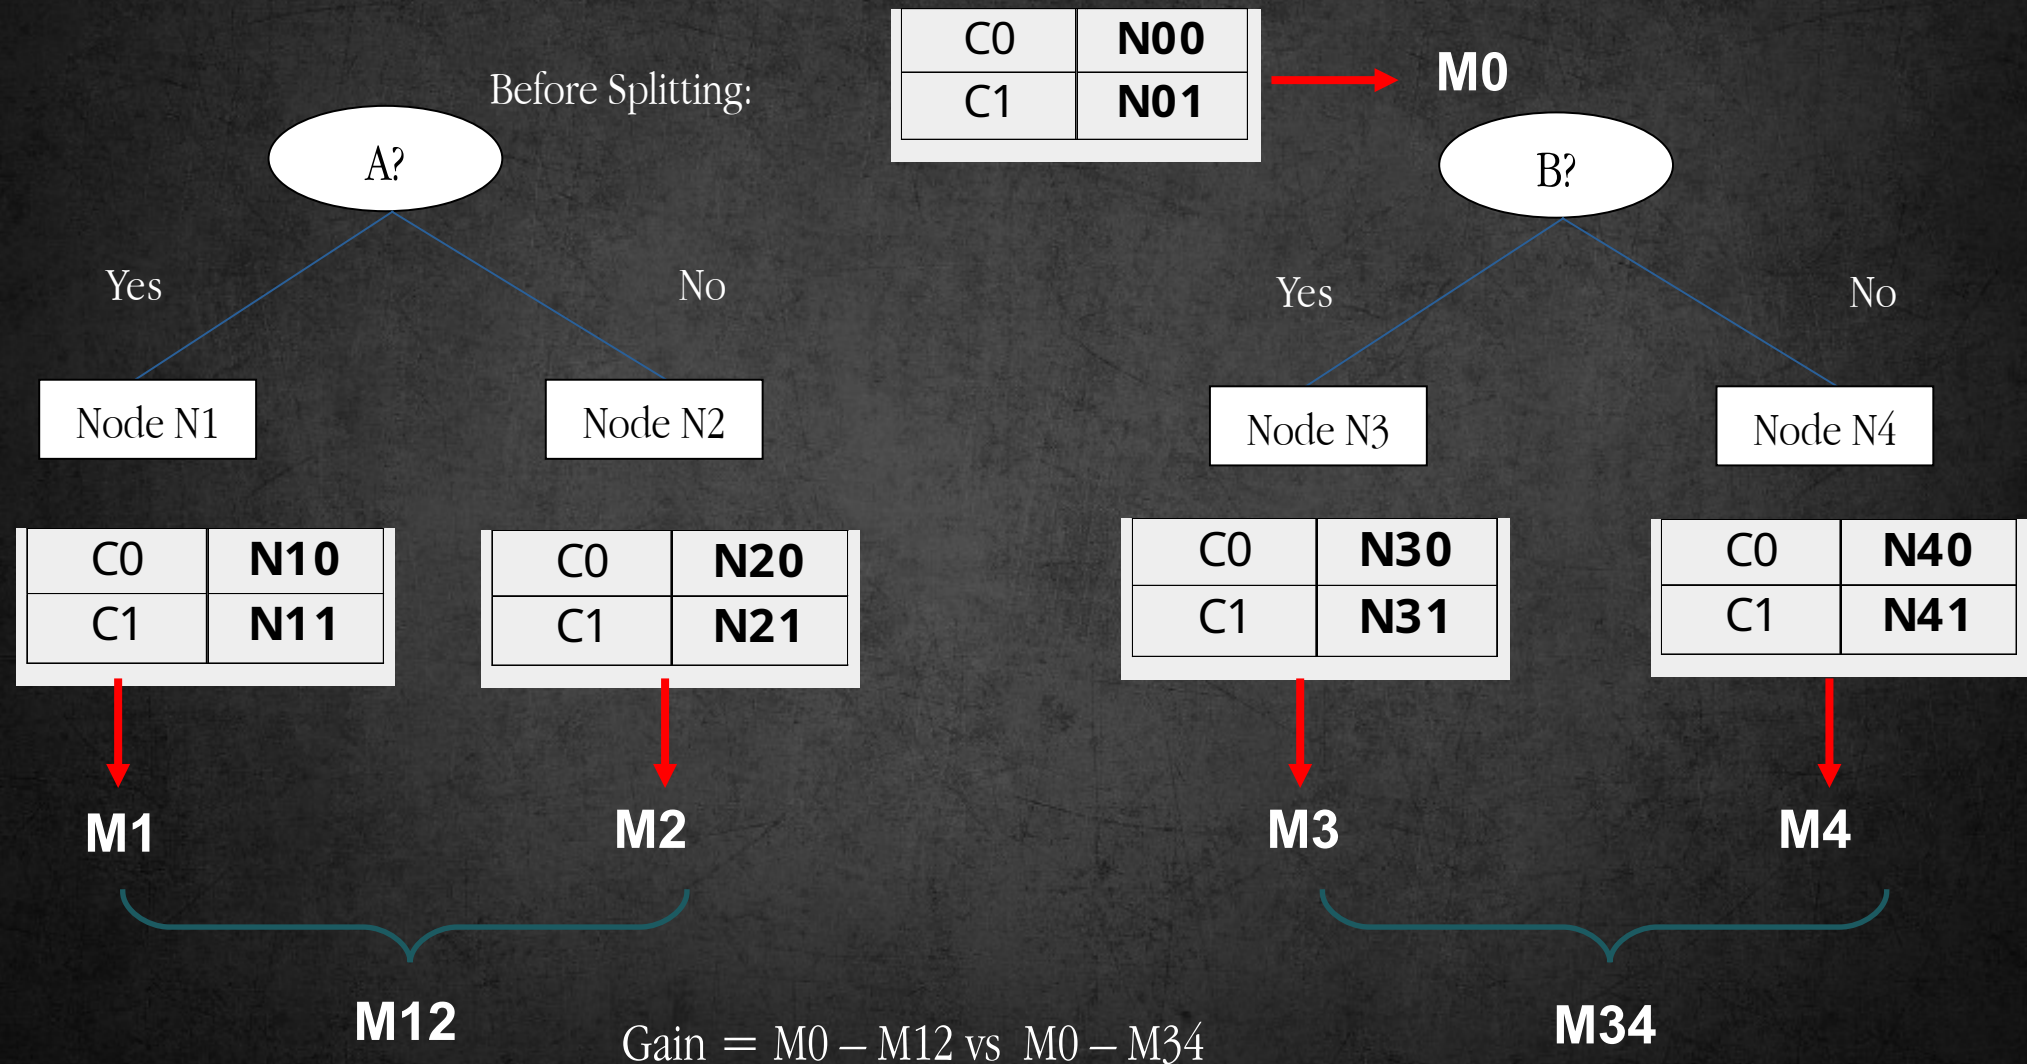
C0: 9
C1: 1

Homogeneous,
Low degree of impurity

Measures of Node Impurity

- Gini Index
- Entropy
- Misclassification error

How to Find the Best Split



Measure of Impurity: GINI

- Gini Index for a given node t :

$$GINI(t) = 1 - \sum_j [p(j|t)]^2$$

(NOTE: $p(j|t)$ is the relative frequency of class j at node t).

- Maximum $(1 - 1/n_c)$ when records are equally distributed among all classes, implying least interesting information
- Minimum (0.0) when all records belong to one class, implying most interesting information

C1	0
C2	6
Gini= 0.000	

C1	1
C2	5
Gini= 0.278	

C1	2
C2	4
Gini= 0.444	

C1	3
C2	3
Gini= 0.500	

Examples for computing GINI

$$GINI(t) = 1 - \sum_j [p(j|t)]^2$$

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Gini = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = 0$$

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Gini = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

C1	2
C2	4

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$Gini = 1 - (2/6)^2 - (4/6)^2 = 0.444$$

Splitting Based on GINI

- Used in CART, SLIQ, SPRINT.
- When a node p is split into k partitions (children), the quality of split is computed as:

$$GINI_{split} = \sum_{i=1}^k \frac{n_i}{n} GINI(i)$$


n_i = number of records at child i

n = number of records at node p

- Best: **Minimum** value

Binary Attributes: Computing GINI Index

- Splits into two partitions
- Effect of Weighing partitions:
 - Larger and Purer Partitions are sought for.



	Parent
C1	6
C2	6
Gini = 0.500	

$$\begin{aligned}
 \text{Gini}(N1) &= 1 - (5/7)^2 - (2/7)^2 \\
 &= 0.4081
 \end{aligned}$$

$$\begin{aligned}
 \text{Gini}(N2) &= 1 - (1/5)^2 - (4/5)^2 \\
 &= 0.3200
 \end{aligned}$$

	N1	N2
C1	5	1
C2	2	4
Gini = 0.3714		

$$\begin{aligned}
 \text{Gini(Children)} &= 7/12 * 0.4081 + \\
 &\quad 5/12 * 0.3200 \\
 &= 0.3714
 \end{aligned}$$

Categorical Attributes: Computing Gini Index

- For each distinct value, gather counts for each class in the dataset
- Use the count matrix to make decisions

Multi-way split

	CarType		
	Family	Sports	Luxury
C1	1	2	1
C2	4	1	1
Gini	0.393		


Two-way split
(find best partition of values)

	CarType	
	{Sports, Luxury}	{Family}
C1	3	1
C2	2	4
Gini	0.400	

	CarType	
	{Sports}	{Family, Luxury}
C1	2	2
C2	1	5
Gini	0.419	

Continuous Attributes: Computing Gini Index

- Use Binary Decisions based on one value
- Several Choices for the splitting value
 - Number of possible splitting values = Number of distinct values
- Each splitting value has a count matrix associated with it
 - Class counts in each of the partitions, $A < v$ and $A \geq v$
- Simple method to choose best v
 - For each v , scan the database to gather count matrix and compute its Gini index
 - Computationally Inefficient! Repetition of work.



Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

Continuous Attributes: Computing Gini Index

- For efficient computation: for each attribute,
 - Sort the attribute on values
 - Linearly scan these values, each time updating the count matrix and computing gini index
 - Choose the split position that has the least gini index

Cheat																				
	Taxable Income																			
Sorted Values	60		70		75		85		90		95		100		120		125		220	
Split Positions	55		65		72		80		87		92		97		110		122		172	
	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>
Yes	0	3	0	3	0	3	0	3	1	2	2	1	3	0	3	0	3	0	3	0
No	0	7	1	6	2	5	3	4	3	4	3	4	3	4	4	3	5	2	6	1
Gini	0.420		0.400		0.375		0.343		0.417		0.400		<u>0.300</u>		0.343		0.375		0.400	

Alternative Splitting Criteria based on INFO

► Entropy at a given node t :

$$\text{Entropy}(t) = - \sum_j p(j|t) \log p(j|t)$$

(NOTE: $p(j|t)$ is the relative frequency of class j at node t).

- Measures homogeneity of a node.
 - Maximum ($\log n_c$) when records are equally distributed among all classes implying least information
 - Minimum (0.0) when all records belong to one class, implying most information
- Entropy based computations are similar to the GINI index computations
- Best: **Minimum** value

Examples for computing Entropy

$$\text{Entropy}(t) = - \sum_j p(j|t) \log_2 p(j|t)$$

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$\text{Entropy} = -0 \log 0 - 1 \log 1 = -0 - 0 = 0$$

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$\text{Entropy} = - (1/6) \log_2 (1/6) - (5/6) \log_2 (5/6) = 0.65$$

C1	2
C2	4

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$\text{Entropy} = - (2/6) \log_2 (2/6) - (4/6) \log_2 (4/6) = 0.92$$

Splitting Based on information gain

➤ Information Gain:

$$GAIN_{split} = Entropy(p) - \left(\sum_{i=1}^k \frac{n_i}{n} Entropy(i) \right)$$

Parent Node, p is split into k partitions;

n_i is number of records in partition i

- Measures Reduction in Entropy achieved because of the split. Choose the split that achieves most reduction (maximizes GAIN)
- Used in ID3 and C4.5
- Disadvantage: Tends to prefer splits that result in large number of partitions, each being small but pure.

Splitting Based on gain ratio

➤ Gain Ratio:

$$GainRATIO_{split} = \frac{GAIN_{Split}}{SplitINFO}$$

$$SplitINFO = - \sum_{i=1}^k \frac{n_i}{n} \log \frac{n_i}{n}$$

Parent Node, p is split into k partitions
 n_i is the number of records in partition i

- Adjusts Information Gain by the entropy of the partitioning (SplitINFO). Higher entropy partitioning (large number of small partitions) is penalized!
- Used in C4.5
- Designed to overcome the disadvantage of Information Gain

Splitting Criteria based on Classification Error

- Classification error at a node t :

$$Error(t) = 1 - \max_i P(i|t)$$

- Measures misclassification error made by a node.
 - Maximum $(1 - 1/n_c)$ when records are equally distributed among all classes, implying least interesting information
 - Minimum (0.0) when all records belong to one class, implying most interesting information

Examples for computing Error

$$Error(t) = 1 - \max_i P(i|t)$$

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Error = 1 - \max(0, 1) = 1 - 1 = 0$$

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Error = 1 - \max(1/6, 5/6) = 1 - 5/6 = 1/6$$


C1	2
C2	4

$$P(C1) = 2/6 \quad P(C2) = 4/6$$


$$Error = 1 - \max(2/6, 4/6) = 1 - 4/6 = 1/3$$

Comparison among Splitting Criteria

- For a two-class problem



Misclassification Error vs. Gini



	Parent
C1	7
C2	3
Gini = 0.42	

$$\begin{aligned}
 \text{Gini}(N1) &= 1 - (3/3)^2 - (0/3)^2 \\
 &= 0.0000
 \end{aligned}$$

$$\begin{aligned}
 \text{Gini}(N2) &= 1 - (4/7)^2 - (3/7)^2 \\
 &= 0.4898
 \end{aligned}$$

	N1	N2
C1	3	4
C2	0	3
Gini = 0.3427		

$$\begin{aligned}
 \text{Gini(Children)} &= 3/10 * 0.000 \\
 &+ 7/10 * 0.4898 \\
 &= 0.3427
 \end{aligned}$$

Gini improves !!

Tree Induction

- Greedy strategy.
 - ⦿ Split the records based on an attribute test that optimizes certain criterion.
- Issues
 - ⦿ Determine how to split the records
 - ⦿ How to specify the attribute test condition?
 - ⦿ How to determine the best split?
 - ⦿ Determine when to stop splitting

Stopping Criteria for Tree Induction

- Stop expanding a node when all the records belong to the same class
- Stop expanding a node when all the records have similar attribute values
- Early termination (to be discussed later)

Decision Tree Based Classification

➤ Advantages:

- Inexpensive to construct
- Extremely fast at classifying unknown records
- Easy to interpret for small-sized trees
- Accuracy is comparable to other classification techniques for many simple data sets
- Unstable


Example: C4.5

- Simple depth-first construction.
- Uses Information Gain
- Sorts Continuous Attributes at each node.
- Needs entire data to fit in memory.
- Unsuitable for Large Datasets.
 - Needs out-of-core sorting.
- You can download the software from:
<http://www.cse.unsw.edu.au/~quinlan/c4.5r8.tar.gz>


Practical Issues of Classification

- Underfitting and Overfitting
- Missing Values
- Costs of Classification
 - See later


Underfitting and Overfitting



Under-fitting and over-fitting




Underfitting and Overfitting




Underfitting: when model is too simple, both training and test errors are large

Overfitting due to Noise



Decision boundary is distorted by noise point

Overfitting due to Insufficient Examples



- Lack of data points in the lower half of the diagram makes it difficult to predict correctly the class labels of that region
- Insufficient number of training records in the region causes the decision tree to predict the test examples using other training records that are irrelevant to the classification task

Notes on Overfitting

- Overfitting results in decision trees that are more complex than necessary
- Training error no longer provides a good estimate of how well the tree will perform on previously unseen records
- Need new ways for estimating errors

Overfitting

➤ Curse of overfitting:

- Related to learning
- Worse when you learning algorithm is better
- No matter how hard you try, it's worse

Overfitting: How to Estimate Generalization Errors

- Re-substitution errors: error on training ($\Sigma e(t)$)
- Generalization errors: error on testing ($\Sigma e'(t)$)
- Methods for estimating generalization errors:
 - Optimistic approach: $e'(t) = e(t)$
 - Pessimistic approach (adds model complexity):
 - For each leaf node: $e'(t) = (e(t) + 0.5)$
 - Total errors: $e'(T) = e(T) + N \times 0.5$ (N: number of leaf nodes)
 - For a tree with 30 leaf nodes and 10 errors on training (out of 1000 instances):
Training error = $10/1000 = 1\%$
Generalization error = $(10 + 30 \times 0.5)/1000 = 2.5\%$
 - Reduced error pruning (REP):
 - Uses validation data set to estimate generalization error

Occam's Razor

- Given two models of similar generalization errors, one should prefer the simpler model over the more complex model
- For complex models, there is a greater chance that it was fitted accidentally by errors in data
- Therefore, one should include model complexity when evaluating a model
- Problem: Not easy to know which is the simpler model

How to Address Overfitting

➤ Pre-Pruning (Early Stopping Rule)

- Stop the algorithm before it becomes a fully-grown tree
- Typical stopping conditions for a node:
 - Stop if all instances belong to the same class
 - Stop if all the attribute values are the same
- More restrictive conditions:
 - Stop if number of instances is less than some user-specified threshold
 - Stop if class distribution of instances are independent of the available features (e.g., using χ^2 test)
 - Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

How to Address Overfitting

➤ Post-pruning

- Grow decision tree to its entirety
- Trim the nodes of the decision tree in a bottom-up fashion
- If validation error improves after trimming, replace sub-tree by a leaf node.
- Class label of leaf node is determined from majority class of instances in the sub-tree
- Can use MDL for post-pruning

Example of Post-Pruning

Class = Yes	20
Class = No	10
Error = 10/30	


Training Error (Before splitting) = 10/30

Pessimistic error = $(10 + 0.5)/30 = 10.5/30$

Training Error (After splitting) = 9/30

Pessimistic error (After splitting) = $(9 + 4 \times 0.5)/30 = 11/30$

PRUNE!



Class = Yes	8
Class = No	4

Class = Yes	3
Class = No	4


Class = Yes	4
Class = No	1

Class = Yes	5
Class = No	1


Examples of Post-pruning

- Optimistic error?
 - Don't prune for both cases
- Pessimistic error?
 - Don't prune case 1, prune case 2
- Reduced error pruning?
 - Depends on validation set

Case 1:



Case 2:



Handling Missing Attribute Values

- If missing values are present, what should I do?
- Missing values affect decision tree construction in three different ways:
 - ⦿ Affects how impurity measures are computed
 - ⦿ Affects how to distribute instance with missing value to child nodes
 - ⦿ Affects how a test instance with missing value is classified

Missing values: Training stage

- Ignoring Data with Missing Attribute Values
 - Do not consider instances and/or attributes
- Most Common Attribute Value
 - The most commonly occurred attribute value is imputed in place of missing values
 - Can be applied within the class
- Distribute the instances
- Method of assigning **ALL** possible Values
 - Multiple training records are created each with possible attribute value in case of missing value: **Noise**
- Null Value Strategy
 - Missing values are treated as a regular value 'Null'
- Imputation methods
 - Prediction model for missing values
 - Imputation with *K*-Nearest neighbors
 - Expectation-maximization algorithm

Computing Impurity Measure: Ignore

Tid	Refund	Marital Status	Taxable Income	Class
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	?	Single	90K	Yes

Missing value

Before Splitting:

$$\text{Entropy}(\text{Parent}) = -0.3 \log(0.3) - (0.7) \log(0.7) = 0.8813$$

	Class = Yes	Class = No
Refund= Yes	0	3
Refund= No	2	4
Refund= ?	1	0

Split on Refund:

$$\text{Entropy}(\text{Refund} = \text{Yes}) = 0$$

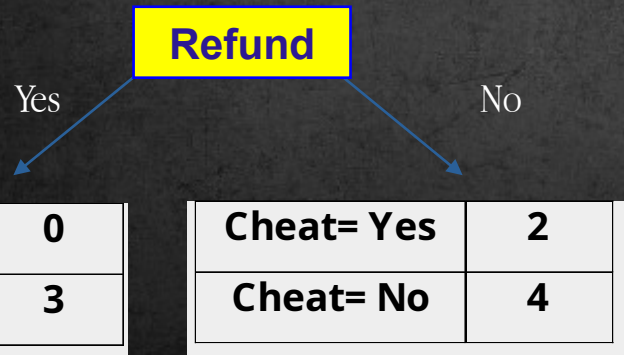
$$\begin{aligned} \text{Entropy}(\text{Refund} = \text{No}) &= -(2/6) \log(2/6) - (4/6) \log(4/6) = 0.9183 \end{aligned}$$

$$\begin{aligned} \text{Entropy}(\text{Children}) &= 0.3 (0) + 0.6 (0.9183) = 0.551 \end{aligned}$$

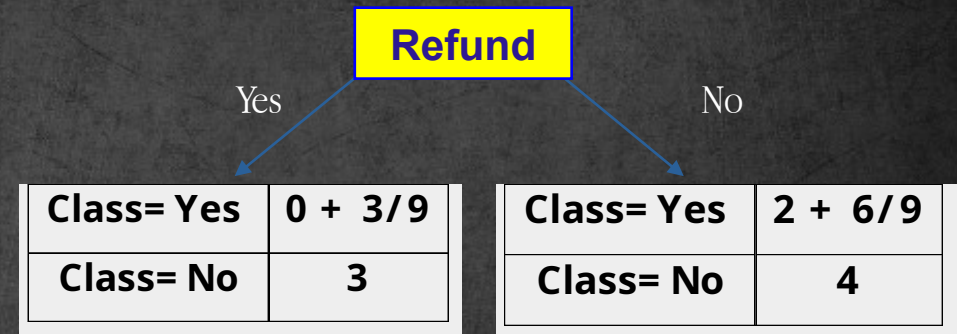
$$\text{Gain} = 0.9 \times (0.8813 - 0.551) = 0.3303$$

Computing Impurity Measure: Distribute Instances

Tid	Refund	Marital Status	Taxable Income	Class
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No



Tid	Refund	Marital Status	Taxable Income	Class
10	?	Single	90K	Yes



Probability that Refund=Yes is $3/9$

Probability that Refund=No is $6/9$

Assign record to the left child with weight = $3/9$
and to the right child with weight = $6/9$

Missing values: Testing stage

➤ Discard Testing Instance

- It may not be acceptable to reject cases for prediction.

➤ Imputation

- Use an imputation method as in training

➤ Null Strategy

- In Null Value strategy as discussed in training time, 'Null' is considered as a special value both at training and testing time.
- Problematic as testing time

C4.5 strategy

- C4.5 does not replace the missing values.
- At the time of selection an attribute at splitting **all instances with known value** of that attribute are used for information gain calculation.
- After selection of an attribute instances with unknown attribute value are split **proportionately** to the split off known values.
- At the time of testing, a test instance with missing value is split into branches according to the portions of training examples falling into those branches.
- This has an advantage over the methods of discarding all incomplete instances in that fewer instances are being discarded.

Classify Instances

New record:

Tid	Refund	Marital Status	Taxable Income	Class
11	No	?	85K	?



	Married	Single	Divorced	Total
Class=No	3	1	0	4
Class=Yes	6/9	1	1	2.67
Total	3.67	2	1	6.67

Probability that Marital Status = Married is $3.67/6.67$

Probability that Marital Status = {Single, Divorced} is $3/6.67$

Other Issues

- Data Fragmentation
- Search Strategy
- Expressiveness

Data Fragmentation

- Number of instances gets smaller as you traverse down the tree
- Number of instances at the leaf nodes could be too small to make any statistically significant decision



Search Strategy

- Finding an optimal decision tree is NP-hard
- The algorithm presented so far uses a greedy, top-down, recursive partitioning strategy to induce a reasonable solution
- Other strategies?
 - Bottom-up
 - Bi-directional

Expressiveness

- Decision tree provides expressive representation for learning discrete-valued function
 - But they do not generalize well to certain types of Boolean functions
 - Example: parity function:
 - Class = 1 if there is an even number of Boolean attributes with truth value = True
 - Class = 0 if there is an odd number of Boolean attributes with truth value = True
 - For accurate modeling, must have a complete tree
- Not expressive enough for modeling continuous variables
 - Particularly when test condition involves only a single attribute at-a-time



Decision Boundary



Border line between two neighboring regions of different classes is known as decision boundary


Decision boundary is parallel to axes because test condition involves a single attribute at-a-time

Oblique Decision Trees



- Test condition may involve multiple attributes
- More expressive representation
- Finding optimal test condition is computationally expensive


Learning Curve



- Learning curve shows how accuracy changes with varying sample size
- Requires a sampling schedule for creating learning curve:
 - Arithmetic sampling (Langley et al.)
 - Geometric sampling (Provost et al.)
- Effect of small sample size:
 - Bias in the estimate
 - Variance of estimate

Bias/Variance of error

- Error = Irreducible error + bias + variance
- Bias
 - Error of the central tendency of the model
- Variance
 - Error of the separation from the central tendency of the model




ROC (Receiver Operating Characteristic)

- Developed in 1950s for signal detection theory to analyze noisy signals
 - Characterize the trade-off between positive hits and false alarms
- ROC curve plots TP rate (S_n) on the y-axis against FP rate ($1-S_p$) on the x-axis
- Performance of each classifier represented as a point on the ROC curve
 - Changing the threshold of algorithm, sample distribution or cost matrix changes the location of the point

ROC Curve


1-dimensional data set containing 2 classes (positive and negative)

any points located at $x > t$ is classified as positive




At threshold t :

TP=0.5, FN=0.5, FP=0.12, FN=0.88




ROC Curve

- (TP rate = TP/P , FP rate = FP/N):
- $(0,0)$: declare everything to be negative class
- $(1,1)$: declare everything to be positive class
- $(1,0)$: ideal
- Diagonal line:
 - Random guessing
 - Below diagonal line:
 - prediction is opposite of the true class



Using ROC for Model Comparison



- No model consistently outperform the other
 - M_1 is better for small FPR
 - M_2 is better for large FPR
- Area Under the ROC curve
- Ideal:
 - Area = 1
- Random guess:
 - Area = 0.5

How to Construct an ROC curve

Instance	$P(+ A)$	True Class
1	0.95	+
2	0.93	+
3	0.87	-
4	0.85	-
5	0.85	-
6	0.85	+
7	0.76	-
8	0.53	+
9	0.43	-
10	0.25	+

- Use classifier that produces posterior probability for each test instance $P(+|A)$
- Sort the instances according to $P(+|A)$ in decreasing order
- Apply threshold at each unique value of $P(+|A)$
- Count the number of TP, FP, TN, FN at each threshold
- TP rate, $TPR = TP/(TP+FN)$
- FP rate, $FPR = FP/(FP + TN)$

How to construct an ROC curve

Threshold \geq

Class	+	-	+	-	-	-	+	-	+	+	
	0.25	0.43	0.53	0.76	0.85	0.85	0.85	0.87	0.93	0.95	1.00
TP	5	4	4	3	3	3	3	2	2	1	0
FP	5	5	4	4	3	2	1	1	0	0	0
TN	0	0	1	1	2	3	4	4	5	5	5
FN	0	1	1	2	2	2	2	3	3	4	5
TPR	1	0.8	0.8	0.6	0.6	0.6	0.6	0.4	0.4	0.2	0
FPR	1	1	0.8	0.8	0.6	0.4	0.2	0.2	0	0	0



ROC Curve:

